

FABRIZIO RUSSO

Research Associate, Imperial College London

 briziorusso.github.io

 [fabriziorusso](https://in.fabriziorusso)

 briziorusso

 Scholar

 dblp

SUMMARY

Research Associate at Imperial College London, starting an AI in Policy Fellowship this coming July, with a PhD in Safe & Trusted AI, developing explainable, interactive causal discovery methods for high-stakes decision-making. As a key contributor to an ERC-POC grant built on my thesis, I'm now working towards applying my research in finance, healthcare and policymaking. I've published at UAI, AAAI, KR, ECAI, IJCAI, and CLeaR. With six years' experience consulting and leading data-science teams in banking, I bridge academia, industry and regulatory bodies to drive human-centred AI innovation.

WORK EXPERIENCE

Research Associate

 Full-time: Jan 2025-present

Imperial College London

 UK

- ERC Proof of Concept grant: Causal Argumentative Learning Assistant (CArLA)
- Research in Causal Discovery, Argumentation and Explainability

Graduate Teaching Assistant

 Part-time: Oct 2021 & 2022

Imperial College London

 UK

- Teaching support for *Introduction to Machine Learning* course

Freelance Consultant

 Part-time: Mar 2022 – Sep 2022

Imperial College Consultants (ICON)

 UK & USA

- Research on outlier identification and forecast for UK commercial properties index

Doctoral Researcher

 Part-time: Jan 2021 – Apr 2021

Ernst & Young (in collaboration with STAI CDT)

 UK

- Research on back-door poisoning attacks identification for Large Language Models

Head of Data Science

 Full-time: Feb 2018 – Sep 2020

4most Europe Ltd

 UK & Germany

- Replicate underwriter behaviour using machine learning
- Commercial credit risk scoring using machine learning
- Design, prototype and build an interactive monitoring tool (Tech of the Year - National Credit Awards '22)

Managing Consultant

 Full-time: Jul 2016 – Jan 2018

4most Europe Ltd

 UK & Canada

Credit Risk Consultant

 Full-time: Jul 2015 – Jun 2016

4most Europe Ltd

 UK

Credit Risk Analyst

 Full-time: May 2014 – Jun 2015

General Electric Capital - UK Home Lending

 UK

EDUCATION

PhD in Safe and Trusted Artificial Intelligence

Imperial College London

📅 Full-time: Oct 2020 – Feb 2025

📍 UK

- Supervisor: Prof. Francesca Toni
- Thesis: Causal Discovery for Trustworthy Artificial Intelligence
- Visiting PhD Student at King's College London ([STAI CDT](#))

MSc in Applied Statistics (Distinction)

University of London - Birkbeck College

📅 Part-time: Oct 2016 – Oct 2018

📍 UK

- Dissertation: Macroeconomic Modelling through VAR and VEC

BSc in Economics of Financial Markets and Institutions (1:1)

University of Rome - Tor Vergata

📅 Full-time: Oct 2009 – Oct 2012

📍 Italy

- Dissertation: The Tobin Tax between Theoretical Wherewithal and Practical Feasibility

PUBLICATIONS

👥 Conference Proceedings

- Li, Zihao and **Fabrizio Russo** (2026). "[Leveraging Large Language Models for Causal Discovery: a Constraint-based, Argumentation-driven Approach](#)". In: *Accepted at the 42nd Conference on Uncertainty in Artificial Intelligence, UAI 2026*.
- **Russo, Fabrizio** and Mark Somers (2026). "Operationalising Relative Causal Knowledge: Backbone Identifiability from Private Reports". In: *Accepted at JoeFest: A Workshop in Honor of Joseph Y. Halpern, co-located with KR / FLoC 2026*.
- **Russo, Fabrizio**, Giuseppe Mazzotta, Carmine Dodaro, Francesco Ricca, and Francesca Toni (2026). "Optimal Correction Sets for Argumentative Causal Discovery". In: *Accepted at LINDA 2026: Logical Approaches to Handling Inconsistent Data, co-located with KR / FLoC 2026*.
- Gehlot, Preesha, Anna Rapberger, **Fabrizio Russo**, and Francesca Toni (2026). "[Heterogeneous Graph Neural Networks for Assumption-Based Argumentation](#)". In: *Proceedings of the 40th Annual AAAI Conference on Artificial Intelligence, AAAI-26*.
- – (2025). "[Heterogeneous Graph Neural Networks for Credulous Acceptance of Assumptions in ABA](#)". in: *Proceedings of the Second International Workshop on Argumentation and Applications co-located with KR 2025*.
- Rapberger, Anna, **Fabrizio Russo**, Antonio Rago, and Francesca Toni (2025). "[On Gradual Semantics for Assumption-Based Argumentation](#)". In: *Proceedings of the 22nd International Conference on Principles of Knowledge Representation and Reasoning, KR 2025*.
- **Russo, Fabrizio** and Francesca Toni (2025). "[Shapley-PC: Constraint-based Causal Structure Learning with a Shapley Inspired Framework](#)". In: *Proceedings of the Fourth Conference on Causal Learning and Reasoning, CLeaR 2025*.
- **Russo, Fabrizio**, Anna Rapberger, and Francesca Toni (2024). "[Argumentative Causal Discovery](#)". In: *Proceedings of the 21st International Conference on Principles of Knowledge Representation and Reasoning, KR 2024*.
- Leofante, Francesco, Hamed Ayoobi, Adam Dejl, Gabriel Freedman, Deniz Gorur, Junqi Jiang, Guilherme Paulino-Passos, Antonio Rago, Anna Rapberger, **Fabrizio Russo**, Xiang Yin, Dekai Zhang, and Francesca Toni (Aug. 2024). "[Contestable AI Needs Computational Argumentation](#)". In: *Proceedings of the 21st International Conference on Principles of Knowledge Representation and Reasoning, KR 2024*.
- **Russo, Fabrizio** and Francesca Toni (2023). "[Causal Discovery and Knowledge Injection for Contestable Neural Networks](#)". In: *ECAI 2023 - 26th European Conference on Artificial Intelligence. Volume 372. FAIA*.
- **Russo, Fabrizio** (2023). "[Argumentation for Interactive Causal Discovery](#)". In: *Proceedings of the 32nd International Joint Conference on Artificial Intelligence, IJCAI 2023*.
- Rago, Antonio, **Fabrizio Russo**, Emanuele Albini, Pietro Baroni, and Francesca Toni (2021). "[Forging Argumentative Explanations from Causal Models](#)". In: *Proceedings of the 5th Workshop on Advances in Argumentation in Artificial Intelligence 2021 co-located with the 20th International Conference of the Italian Association for Artificial Intelligence (AIxIA 2021)*.

📖 Journal Articles

- Rago, Antonio, **Fabrizio Russo**, Emanuele Albini, Pietro Baroni, and Francesca Toni (2023). "[Explaining Classifiers' Output with Causal Models and Argumentation](#)". In: *Journal of Applied Logics – IfCoLog Journal of Logics and their Applications, Special Issue: Advances in Argumentation in AI 10.3*, pages 421–449.

SERVICE

Teaching

- **COMP70076 Ethics, Fairness and Explanation in AI:** Guest Lecture, Imperial College London (Spring 2026)
- **Data Science In Action:** Guest Lecture, Luiss Guido Carli, (Autumn 2025)
- **COMP70050 Introduction to Machine Learning:** Tutorials and Marking, Imperial College London (2021, 2022)
- **MSc in Fintech - Politecnico di Milano & UCL:** Guest Lecture, UCL (Spring 2022)

Supervision

- Dhruv Himatsingka (MEng 2026): Bridging Gradual Semantics and Graph Neural Networks for Argumentative Causal Discovery
- Zihao Li (MSc 2025): Leveraging Large Language Models for Causal Discovery
- Yeva Hunanyan (MSc 2025): Exploring Semantics for Argumentative Causal Discovery
- Preesha Gehlot (MEng 2025): Graph Neural Networks for Approximate Argumentative Reasoning

Reviewing

- **Artificial Intelligence Journal:** Reviewer (2025, 2026)
- **KR:** Reviewer (2026)
- **IJCAI-ECAI:** Reviewer (2023, 2026)
- **AAAI:** Reviewer (2026)
- **NeurIPS:** Reviewer (2024, 2025)
- **ICLR:** Reviewer (2025)
- **ICML:** Reviewer (2022)
- **AISTATS:** Reviewer (2022, 2024); Top-Reviewer (2023)
- **AAMAS:** Reviewer (2023)

Organising

- **XAI Seminar Series @ Imperial:** Co-organiser (2021–present)
- **XAI-FIN at ICAIF:** Co-Chair (2025); Assistant Chair (2023, 2024); Workshop co-organiser (2022); Reviewer (2021)
- **KCL-ICL Workshop on Causality:** Co-organiser (2025)

Invited Talks

- Minds and Machines Workshop, University of Queensland (2025)
- Lendable Journal Club (2025)
- Lloyds Banking Group Summer School (2025)
- STAI CDT Student Seminar Series (2023)
- Danmarks Nationalbank Conference (2019)