

FABRIZIO RUSSO

AI in Policy Fellow, Imperial College London

@ brizio.russo@gmail.com

briziorusso.github.io

in fabriziorusso

briziorusso

Scholar

dblp

SUMMARY

With six years' experience consulting and leading data-science teams in banking, I bridge academia, industry and regulatory bodies to drive human-centred AI innovation. Building on my PhD thesis in Safe & Trusted AI, and my postdoctoral research on explainable and contestable causal discovery (CArLA), I am now an AI in Policy Fellow at Imperial, translating explainable, interactive causal discovery research into practical tools for finance, healthcare and policymaking. I've published at established peer-reviewed venues while building an independent academic profile aimed at, and grounded on, real-world impact.

WORK EXPERIENCE

Imperial College London

- + AI in Policy Fellow 📅 Full-time: Sep 2026-present 📍 UK
- + Research Associate 📅 Full-time: Mar 2025-Aug 2026 📍 UK
 - ERC Proof of Concept grant: Causal Argumentative Learning Assistant (CArLA)
- + Research Assistant 📅 Full-time: Jan 2025-Feb 2025 📍 UK
- + Graduate Teaching Assistant 📅 Part-time: Oct 2021 & 2022 📍 UK

Consultant

- + Imperial College Consultants (ICON) 📅 Part-time: Mar 2022 – Sep 2022 📍 UK & USA
 - Research on outlier identification and forecast for UK commercial properties index
- + Ernst & Young (in collaboration with STAI CDT) 📅 Part-time: Jan 2021 – Apr 2021 📍 UK
 - Research on back-door poisoning attacks identification for Large Language Models

4most Europe Ltd

- + Head of Data Science 📅 Feb 2018 – Sep 2020 📍 UK & Germany
 - Responsible for the development and delivery of data science projects for financial institutions
 - Design, prototype and build an interactive ML model monitoring tool (Tech of the Year - National Credit Awards '22)
- + Managing Consultant 📅 Jul 2016 – Jan 2018 📍 UK & Canada
- + Credit Risk Consultant 📅 Jul 2015 – Jun 2016 📍 UK

Credit Risk Analyst

📅 Full-time: May 2014 – Jun 2015

General Electric Capital - UK Home Lending

📍 UK

EDUCATION

PhD in Safe and Trusted Artificial Intelligence

📅 Full-time: Oct 2020 – Feb 2025

Imperial College London

📍 UK

- Supervisor: Prof. Francesca Toni
- Thesis: Causal Discovery for Trustworthy Artificial Intelligence
- Visiting PhD Student at King's College London (STAI CDT)

MSc in Applied Statistics (Distinction)

📅 Part-time: Oct 2016 – Oct 2018

University of London - Birkbeck College

📍 UK

- Dissertation: Macroeconomic Modelling through VAR and VEC

BSc in Economics of Financial Markets and Institutions (1:1)

📅 Full-time: Oct 2009 – Oct 2012

University of Rome - Tor Vergata

📍 Italy & Spain

- Dissertation: The Tobin Tax between Theoretical Wherewithal and Practical Feasibility
- Exchange Program: University of Madrid - Autonoma (Oct 2011 – Jun 2012)

PUBLICATIONS

- Li Z, Russo F. [Leveraging Large Language Models for Causal Discovery: a Constraint-based, Argumentation-driven Approach](#). Proceedings of the 42nd Conference on Uncertainty in Artificial Intelligence. Pages 3631–3667. 2026.
- Gehlot P, Rapberger A, Russo F, Toni F. [Heterogeneous Graph Neural Networks for Assumption-Based Argumentation](#). Proceedings of the 40th AAAI Conference on Artificial Intelligence. Pages 19117–19125. 2026.
- Rapberger A, Russo F, Rago A, Toni F. [On Gradual Semantics for Assumption-Based Argumentation](#). Proceedings of the 22nd International Conference on Principles of Knowledge Representation and Reasoning. Pages 512–522. 2025.
- Russo F, Toni F. [Shapley-PC: Constraint-based Causal Structure Learning with a Shapley Inspired Framework](#). Proceedings of the 4th Conference on Causal Learning and Reasoning. Pages 292–339. 2025.
- Russo F, Rapberger A, Toni F. [Argumentative Causal Discovery](#). Proceedings of the 21st International Conference on Principles of Knowledge Representation and Reasoning. Pages 938–949. 2024.
- Leofante F, Ayoobi H, Dejl A, Freedman G, Gorur D, Jiang J, Paulino-Passos G, Rago A, Rapberger A, Russo F, Yin X, Zhang D, Toni F. [Contestable AI Needs Computational Argumentation](#). Proceedings of the 21st International Conference on Principles of Knowledge Representation and Reasoning. Pages 888–896. 2024.
- Russo F, Toni F. [Causal Discovery and Knowledge Injection for Contestable Neural Networks](#). Proceedings of the 26th European Conference on Artificial Intelligence. Pages 2025–2032. 2023.
- Russo F. [Argumentation for Interactive Causal Discovery](#). Proceedings of the 32nd International Joint Conference on Artificial Intelligence. Pages 7091–7092. 2023.
- Rago A, Russo F, Albini E, Toni F, Baroni P. [Explaining Classifiers' Outputs with Causal Models and Argumentation](#). Journal of Applied Logics – IfCoLog Journal of Logics and their Applications. Volume 10, Number 3. Pages 421–449. 2023.

SERVICE

Teaching

- **COMP70076 Ethics, Fairness and Explanation in AI:** Guest Lecture, *Causal Discovery for Trustworthy AI*, Imperial College London, London UK (Spring 2026)
- **Bachelor Degree in Data Science In Action:** Guest Lecture, *Causal Discovery for Trustworthy AI*, Luiss Guido Carli, Rome Italy (Autumn 2025)
- **MSc in Fintech - Politecnico di Milano & UCL:** Guest Lecture, *From Credit Risk to Explainable AI Research*, UCL, London UK (Spring 2022)
- **COMP70050 Introduction to Machine Learning:** Tutorials and Marking, Imperial College London (2021, 2022)

Supervision

- Sam Waugh (MSc 2026): Causal ABA Learning
- Tiancheng Yang (MSc 2026): Explainable AI for Answer Sets and its application to Causal Discovery
- Dhruv Himatsingka (MEng 2026): Bridging Gradual Semantics and Graph Neural Networks for Argumentative Causal Discovery
- Zihao Li (MSc 2025): Leveraging Large Language Models for Causal Discovery
- Yeva Hunanyan (MSc 2025): Exploring Semantics for Argumentative Causal Discovery
- Preesha Gehlot (MEng 2025): Graph Neural Networks for Approximate Argumentative Reasoning

Reviewing

- **Artificial Intelligence Journal:** Reviewer (2026, 2025)
- **KR:** Reviewer (2026)
- **IJCAI-ECAI:** Reviewer (2026, 2023)
- **AAAI:** Reviewer (2026)
- **NeurIPS:** Reviewer (2025, 2024)
- **ICLR:** Reviewer (2025)
- **AISTATS:** Reviewer (2024, 2022); Top-Reviewer (2023)
- **AAMAS:** Reviewer (2023)
- **ICML:** Reviewer (2022)

Organising

- **XAI Seminar Series @ Imperial:** Co-organiser (2021–present). Grown from 40 to 400+ subscribers.
- **XAI-FIN at ICAIF:** Co-Chair (2025); Assistant Chair (2023, 2024); Workshop co-organiser (2022). 100+ attendees.
- **KCL-ICL Workshop on Causality:** Co-organiser (2025)

Invited Talks

- *Causal AI You Can Challenge: From Data to Decisions in Banking*. Lloyds Banking Group Summer School, London UK (2026)
- *CausAlly: Collaborative Causal Discovery for AI-Aided Decision-Making*. Minds and Machines Workshop, University of Queensland, Brisbane Australia (2025)
- *Causal Discovery for Trustworthy AI*. Lendable Journal Club, London UK (2025)
- *Causal Discovery for Trustworthy AI*. Lloyds Banking Group Summer School, London UK (2025)
- *Causal Discovery for Contestable and Explainable AI: From Contestable Neural Networks to Transparent Causal Discovery*. Robert Gordon University, Aberdeen UK (2024)
- *Argumentative Causal Discovery*. STAI CDT Student Seminar Series, London UK (2023)
- *Credit Risk Modelling: Data and Techniques Used in the UK Banking Industry*. Danmarks Nationalbank Conference, Copenhagen Denmark (2019)